

2025-2026 Master Thesis Topics in OpenMP and ML/AI

John Burns, ADA University

July 16, 2025

Overview

- 1 Introduction
- 2 Topic 1: Optimizing Deep Neural Network Training with OpenMP
- 3 Topic 2: Help! Advising OpenMP Parallelization
- 4 Topic 3: Parallel Graph Neural Networks with OpenMP
- 5 Topic 4: Energy-Efficient ML Model Training with OpenMP
- 6 Topic 5: OpenMP Enhancement of R Selected Contributed Packages

About myself and my research interests

- My PhD was in the field of large-scale (at the time!) modelling and simulation
- From my industry background I also have an interest in performance, reliability and scalability
- From my teaching I have an interest in Computer Architecture and Distributed Systems
- I welcome any suggestions you may have in the intersection of these fields
- I am very interested in OpenMP projects - from code and performance analysis through to application development
- As Energy Efficiency and Optimization is so important now, why not use ARM64 cloud platforms for these research topics.

Topic 1: Optimizing Deep Neural Network Training with OpenMP

- **Objective:** Enhance the performance of deep neural network (DNN) training by parallelizing compute-intensive operations (e.g., matrix multiplications, gradient computations) using OpenMP on multi-core CPUs.
- **Description:** Implement and evaluate OpenMP-based parallelization strategies for DNN training in frameworks like PyTorch or TensorFlow. Focus on optimizing data parallelism and model parallelism for large-scale datasets.
- **Relevance:** Reduces training time for DNNs, making them more feasible for resource-constrained environments.
- **Methodology:** Develop parallelized versions of backpropagation and optimization algorithms, benchmark performance on multi-core systems, and compare with GPU-based implementations.

Topic 1: References

- **Rakhimov, M., et al.**, "Parallel Approaches in Deep Learning: Use Parallel Computing," *International Conference on Future Networks and Distributed Systems (ICFNDS)*, 2023.
<https://doi.org/10.1145/3644713.3644738>
- **Zhang, Y., et al.**, "Exploiting Parallelism Opportunities with Deep Learning Frameworks," *ACM Transactions on Architecture and Code Optimization*, vol. 18, no. 1, 2020.
<https://doi.org/10.1145/3431388>

Topic 2: Help! Advising OpenMP Parallelization

- **Objective:** Reduce the programming burden of OpenMP directive annotation tasks.
- **Description:** Research DeepTyper-based learning architecture for advising OpenMP Annotation / Parallelization.
- **Relevance:** There are over 30 distinct OpenMP directives in version 5.1 of the specifications including **Parallel, Workshareing, Synchronization, Tasks, Device/Offloading** (for accelerators like GPUs) constructs.
- **Methodology:** Using CodeT5+ and/or DeepTyper approaches for improved directive selection.

Topic 2: References

- **Pornmaneerattanatri, S., et al.,** "Automatic Parallelization with CodeT5+: A Model for Generating OpenMP Directives," *International Workshop on Large Language Models and HPC*, 2024.
- **Shen, Y., et al.,** "A machine learning method to variable classification in OpenMP," *Concurrency and Computation: Practice and Experience*, 2023. <https://doi.org/10.1002/cpe.7746>
- **Kadosh, T., et al.,** "Advising OpenMP Parallelization via A Graph-Based Approach with Transformers" *OpenMP: Advanced Task-Based, Device and Compiler Programming* (pp.3-17)

Topic 3: Parallel Graph Neural Networks with OpenMP

- **Objective:** Enhance the scalability of graph neural networks (GNNs) for large-scale graph data using OpenMP parallelization.
- **Description:** Implement OpenMP-based parallel processing for GNN operations like message passing and aggregation, targeting applications in social networks or bioinformatics.
- **Relevance:** GNNs are computationally expensive for large graphs; OpenMP can improve training and inference speed on multi-core systems.
- **Methodology:** Develop parallel GNN algorithms, test on datasets like OGB (Open Graph Benchmark), and compare performance with serial implementations.

Topic 3: References

- **Meng, Z., et al.**, "OpenMP Parallelization and Optimization of Graph-Based Machine Learning Algorithms," *International Workshop on OpenMP (IWOMP)*, 2016
- **Zhou, J., et al.**, "Graph Neural Networks: A Review of Methods and Applications," *AI Open*, vol. 1, 2020.
<https://doi.org/10.1016/j.aiopen.2021.01.001>

Topic 4: Energy-Efficient ML Model Training with OpenMP

- **Objective:** Develop energy-efficient ML training pipelines using OpenMP to optimize resource utilization on multi-core CPUs.
- **Description:** Investigate OpenMP-based parallelization to reduce energy consumption in ML tasks, focusing on dynamic thread management and workload balancing.
- **Relevance:** Energy efficiency is critical for sustainable AI; OpenMP can optimize CPU-based training for green computing.
- **Methodology:** Implement energy-aware parallel algorithms, measure energy consumption using PowerPoint, and compare with traditional ML training pipelines.

Topic 4: References

- **Garcia, A., et al.**, "DNN Is Not All You Need: Parallelizing Non-Neural ML Algorithms on Ultra-Low-Power IoT Processors," *ACM Transactions on Embedded Computing Systems*, 2023.
<https://doi.org/10.1145/3570152>
- **Strubell, E., et al.**, "Energy and Policy Considerations for Deep Learning," *ACM Computing Surveys*, vol. 53, no. 3, 2020.
<https://doi.org/10.1145/3372822>

Topic 5: OpenMP Enhancement of R Selected Contributed Packages

- **Objective:** Enhance the performance of selected package using OpenMP.
- **Description:** Today there are over 22000 Contributed Packages in cran-r.
- **Relevance:** R is one of the most popular languages for data science, with 31% of data scientists regularly using it, according to a 2025 source..
- **Methodology:** Select a popular R contrib package, (`install.packages("packageRank")`) or some other package of your choice. Compile locally and test with and without OpenMP directives. Benchmark performance and contribute to the community

Topic 5: References

- As of July 17, 2025, there are approximately 1.1–1.3 million R programmers globally, based on a total R user base of 2.1–2.2 million and a 50–60% programmer ratio.
- Start here: `install.packages("packageRank")`

Conclusion

- These topics combine OpenMP's parallel computing capabilities with cutting-edge ML/AI challenges.
- They address performance, scalability, privacy, and sustainability in AI systems.
- Obviously knowledge/interest in C and Parallel Computing is a requirement
- A good place to start: OpenMP Architecture Review Board, "OpenMP API Specification: Version 5.2," 2024.
<https://www.openmp.org/specifications/>